Pytorch

pytorch□tensorflow□□□□

pytorch□□□□□□□□

- □□□□□□□□□□□□□□□□□
- □□□□□□□□□□□□□□□□□□

tensorflow□□□□□□□□

- □□□□□□□□
- □□□□□□□□□□□□□□□□

datasets

`from torchvision import datasets`

□□□□□□

- `.□□□□□□`
  - `root` □□□□□□□□□□
  - `train` □□□□□□□□□□
  - `download` □□□□□□
  - `transform` □□□□□□□□□□

Dataset

□ `torch.utils.data` □□□□□□□□□□□□□□□□□□□□□□

□□□□□□□□□□□□ `torch.utils.data.Dataset` □□□□□□□□□□□□

- `__len__()` □□□□□□□□□□□□□□□
- `__getitem__(idx)` □□□□□□ `idx` □□□□□□□□□□□□□□□□□□

DataLoader

□ `torch.utils.data` □□□□ `Dataset` □□□□□□□□□□□

- □□□□□□ `batch_size` □□□□□□□□□□□□□□□□□□batch□□
- □□□□□□ `shuffle=True` □□□□ epoch □□□□□□□□□□
- □□□□□□□ `num_workers` □□□□□□□□/□□□□□□□□□

□□	□□
`dataset`	□□□□Dataset□□
`batch_size`	□□□□□□□□□
`shuffle`	□□□□□epoch□□□□
`num_workers`	□□□□□□□□□□□□CPU□□□□
`drop_last`	□□□□□□□□□□□□□□□
`pin_memory`	□□□□□□□□□GPU□□

Transfrom

□ `torchvision` □□□□□□□□□□□

□□	□□
`ToTensor()`	□ PIL □□□□□□□□□□
`Normalize(mean, std)`	□□□□□□□□□□□□□□□□
`RandomCrop(size)`	□□□□
`RandomHorizontalFlip()`	□□□□□□□

□□□□

□ `torch.nn.Function` □ $$ ReLU(x) = max(0,x) $$

- Relu
  - □□□□□□$[0, +\infty)$
  - □□□□□□□□□□Sigmoid□□□
  - □□□□□□□□□□□□

$$ tanh(x) = \frac{e^x - e^{-x}}{e^x+e^{-x}} $$

- Tanh
  - □□□□□□(-1, 1)
  - □□□□□□□□□□□□□□□□□□□□□□□□□□□1□□
  - □□□□□□□□□□□□□□□0□□

$$ \sigma (x) = \frac{1}{1+e^{-x}} $$

- Sigmoid
  - □□□□□□□0,1□
  - □□□□□□□□□□□□□□□□□□□□□□□□□0□□
  - □□□□□□□□□□□□□□□□□□□□□□□□

$$ GELU(x) = x \cdot \Phi(x) $$

- GELU□□□□□□□□□□□□
  - □ ReLU □□□□□□□□□□□□□□□□□□□□□□
  - $\Phi(x)$□□□□□□□□□□□□□□□□□□□
  - □□ Transformer □□

$$ Swish(x) = x \cdot \sigma(\beta x) $$

- Swish
  - σ□Sigmoid□□□β□□□□□□□□□□□□□□□□□□□□□
  - □□□□□□□□□□□□□□□□□□□□□□**ReLU**□
  - □□□Transformer□□□□□□□□□□□□□□□□□□□□□□□

```
class Swish(nn.Module):
    def __init__(self,beta = 1.0):
        super().__init__()
        self.beta = beta

    def forward(self,x):
        return x * torch.sigmoid(self.beta * x)
```

$$ PReLU(x) = \begin{cases} x & \text{if } x \geq 0 \\ ax & \text{if } x < 0 \end{cases} $$

- PReLU
  - □□□□□□□□□□□□□a□□□□□□□
  - □□□□□□CNN□□□□□□□□□□□□□

□□□□

□ `torch.nn` □

1. □□□□□□□□**MSE**□

   - □□□

   $$ MSE = \frac{1}{n}\sum_{i=1}^n(y_i-\hat{y_i})^2 $$

   - □□□□□□□
   - `nn.MSELoss()`

2. □□□□□□□**Cross-Entropy**□

   - □□□

   $$ CE = -\sum_{i=1}^n y_i\log(\hat{y_i}) $$

   - □□□□□□□
   - `nn.CrossEntropyLoss()`

□□□

□ `torch.optim` □

□□□□□□□□□□□□□□□□

1. □□□□□□□ `nn.Linear` □ `nn.Conv2d` □□
2. □□□□□□□□ `optim.SGD` □ `optim.Adam` □□
3. □□□□□□□□□□□□□
   - `optimizer.zero_grad()` □□□□□□□
   - `loss.backward()` □□□□□□□
   - `optimizer.step()` □□□□□□□

SGD□□□□□□□□□

- □□□
  - model.paramaters()□□□□□□
  - lr□□□□□
  - momentum□□□□
  - dampening□□□□□□□□ 0□

- weight_decay□L2 □□□□□□ 0

$$ v_t = \beta v_{t-1} + □1-\beta) \frac{\partial J}{\partial w} \ w = w - \alpha v_t $$

Adam□□□□□□□□

- □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□
- □□□
    - model.paramaters()□□□□□□
    - lr□□□□□
    - betas□□□□
        - β_1□□□□□□□
        - β_2□□□□□□□
    - eps□□□□□□□□□□□0□
    - weight_decay□□□□□□L2□□□□

1. □□□□□□□□□□□□□□□□□□□

g_t□□□ $$ m_t = \beta_1 \cdot m_{t-1} + (1 - \beta *1)\cdot g_t \ v_t = \beta_2\cdot v*{t-1}+(1-\beta_2)\cdot g_t^2 $$

2. □□□□

$$ \hat m_t = \frac{m_t}{1-\beta_1^t}\ \hat v_t = \frac{v_t}{1-\beta_2^t} $$

3. □□□□

$$ \theta_{t+1} = \theta_t - \frac{\alpha \cdot \hat m_t}{\sqrt{\hat v_t + \epsilon}} $$

RMSprop

- □□□□□□□□□□□□□
- □□
    - model.paramaters()□□□□□□
    - lr□□□□□
    - alpha□□□□□□□□□□
    - eps□□□□□□□□□□□0□
    - weight_decay□□□□□□L2□□□□

$$ v_t = \beta v_{t-1} + □1-\beta) (\frac{\partial J}{\partial w})^2 \ w = w - \alpha \frac{\frac{\partial J}{\partial w}}{\sqrt{v_t + \epsilon}} $$

□□□□□□

□□□□□□□□□□

```
from torch.optim import lr_scheduler
```

StepLR□□□□□□□

□□	□□
`optimizer`	□□□□□
`step_size`	□□□□epoch□□□□□□□
`gamma`	□□□□□□□□□0.1□

`last_epoch`	□□□□epoch□□

`scheduler.step()` □□□□□□□**scheduler**□□□□□□□□□

ReduceLROnPlateau (□□□□□□□□)

□□	□□
`optimizer`	□□□□□
`mode`	`'min'`□□□□□□□□□`'max'`□□□□□□□□
`factor`	□□□□□□□
`patience`	□□□□□ epoch □□□□□□□□□□□
`threshold`	□□□□□□□□□□□□□□□□□□□
`threshold_mode`	`'rel'`□□□□□□□ `'abs'`□□□□□□
`cooldown`	□□□□□□□□□□□ epoch □□□□□□□
`min_lr`	□□□□□
`verbose`	□□□□□□□□

`scheduler.step(val_loss)` □□□□□□□

CosineAnnealingLR (□□□□)

□□	□□
`optimizer`	□□□□□
`T_max`	□□□□□□□epoch□□□
`eta_min`	□□□□□
`last_epoch`	□□□□epoch□□

`scheduler.step()` □□□□□□

Conv

□ `torch.nn` □□□□□□□□□□□□□□□□□□□□□□□□□□□

□□□□□□□□□$(in\_channels,H,W)$

- □□□□□□□□□□□□□□□□□□□□□□□□□□
- □□□□□□□□□□□□□□□□□□□□□□

□□□

- `in_channels` □□□□□□□□□□RGB□□□3□
- `out_channels` □□□□□□□□□□□□□
- `kernel_size` □□□□□□□□□□□□
- `stride` □□□□□
- `padding` □□□□□

- `bias` □□□□
- `dilation` □□□□□□□□□□□□□□□□□□□
- `groups` □□□□□□□□□□□□□□□□□□□□□□

□□□□□□□ $$ H_{out} = \frac{H_{in} + 2 \times padding - dilation\times(kernel_size - 1) - 1 }{stride} $$

TransPosed Conv

□ `torch.nn` □□□□□□□□□□□□□□□□□□□□□□□□□ `nn.ConvTranspose2d`

□□□□□□□□□□□

□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□

□□□□

1. □□□□□□□□□□□ `s-1` □□□0
2. □□□□□□□□□□ `k-p-1` □□□0
3. □□□□□□□□□□□□□
4. □□□□□□□

□□□

- `in_channels` □□□□□□□□□□RGB□□3□
- `out_channels` □□□□□□□□□□□□□
- `kernel_size` □□□□□□□□□□□□
- `stride` □□□□□□□0□□□
- `padding` □□□□□
- `bias` □□□□
- `out_padding` □□□□□□□□□□□□□□□□□□□□

image-20250805130358084

image-20250805130421265

□□□□□□□ $$ H_{out} = (H_{in} - 1) \times stride - 2 \times padding + kernel_size+output_padding $$

Pool

□ `torch.nn` □□□□□□□□□□□□□□□□□□□□□□□

□□□□□

- `nn.MaxPool2d` □□□□□□□□□□□□□□□
- `nn.AvgPool2d` □□□□□□□□□□□□□
- `nn.AdaptiveMaxPool2d` □□□□□□□□□□□□□□□□

□□□

- `kernel_size` □□□□□□□
- `stride` □□□□□□□□ `kernel_size` □
- `padding` □□□□□
- `dilation` □□□□□

□□□□□□□□ $$ H_{out} = \frac{H_{in} + 2 \times padding - dilation\times(kernel_size - 1) - 1 }{stride} $$

Linear

□ `torch.nn` □□□□□□□□□□□□□□□□□□□□□□□□□□□□ $$ y = x \cdot W^T + b $$ □□□

- x□□□□□□□□□□□□□$in_features$

□□□

- in_feature:□□□□□□
- out_feature:□□□□□□
- bias:□□

□□□

Flatten

□ `torch.nn` □□□□□□□□□□□□□□□□□□□□

- □□□□□□□□□□□□□□□□□□□ (batch_size, channels, height, width))
- □□□□□□□□ 2D □□□□□□□ (batch_size, channels * height * width)

Normalization

□□ `torch.nn` □

Batch Normalization

□□□□□□□□□□□□□□□□□□□

□□□□□□□□ `nn.BatchNorm1d`

- □□□□□□□□

□□□□□□□ `nn.BatchNorm2d` □□□□□□□□□□□□□□□□Batch□□□□□□□

- □□□□□□□□ `[Batch_size, Channels, Height, Width]`
- □□□□□□□□□□Batch□□□□□□□□□□□□□□□□□□□□□□□□**Batch**□□□
- □□□□□□□□Channels□

□□□

- □□□□□□□□□□□□
- □□□batch□□□□□□□□□□□□□□□□
- □□□□□□□□□□□□□□□□□□□□□

Layer Normalization

□□□□□□□□□□□□□□□□□□□□□□□

`nn.LayerNorm`

□□□

- □□□□□□□□□□□□
- □□**RNN/Transformer**

Group Normalization

`nn.GroupNorm`

[illegible]

[illegible]

- `num_groups` [illegible]
- `num_channels` [illegible]
- `affine` [illegible]

Upsample

`nn.Upsample`

[illegible]

[illegible]

[illegible]	[illegible]	[illegible]
`size`	int[illegible]tuple	[illegible]`(256, 256)`[illegible]None[illegible]`scale_factor`[illegible]
`scale_factor`	float[illegible]tuple	[illegible]`2`[illegible]2[illegible]`(2, 3)`[illegible]2[illegible]3[illegible]
`mode`	str	[illegible]`'nearest'`[illegible]`'bilinear'`[illegible]`'bicubic'`[illegible]`'nearest'`[illegible]
`align_corners`	bool	[illegible]`bilinear/bicubic`[illegible]False[illegible]

Dropout

[illegible] `torch.nn` [illegible]0[illegible]

[illegible] **`.train()`** [illegible] **`.eval()`** [illegible]

Interpolate

[illegible] `torch.nn.functional.interpolate` [illegible]

[illegible]/[illegible]

[illegible]

[illegible]	[illegible]	[illegible]
`size`	`(target_h, target_w)`	[illegible]
`scale_factor`	`float` [illegible] `(float, float)`	[illegible] `2.0` [illegible]2[illegible]
`mode`	`'nearest'`, `'bilinear'`, `'bicubic'`, `'area'`	[illegible]
`align_corners`	`bool`	[illegible]`True` [illegible]

[illegible]

[illegible]

- `torch.utils.checkpoint`
  - [illegible]

1. □□□□□□□□□□□□□□ `checkpoint` □□□
2. □□□□□□□ `checkpoint` □□□□□□□□

- □□
  - `fn` □□□□□□□□□□□
  - `**args` □□□□□□□□□□□□

□□

```
import torch
import torch.nn as nn
from torch.utils.checkpoint import checkpoint

class BigModel(nn.Module):
    def __init__(self):
        super().__init__()
        self.layer1 = nn.Linear(1000, 1000)
        self.layer2 = nn.Linear(1000, 1000)
        self.layer3 = nn.Linear(1000, 1000)

    def forward(self, x):
        # □□□□□□□□□□□□□
        x = checkpoint(self.layer1, x)  # □□□□□□ self.layer1(x)
        x = checkpoint(self.layer2, x)
        x = self.layer3(x)  # □□□□□□□□□□□□□□□□
        return x

## □□
model = BigModel()
inputs = torch.randn(32, 1000)
outputs = model(inputs)
loss = outputs.sum()
loss.backward()  # □□□□□□□□□□layer1□layer2□□□□
```

- `torch.utils.checkpoint_sequential` □**Sequential□□□□**
  - □□
    - `functions` □Sequential□□□□□□□
    - `input` □□□□□□□□□□□□
    - `segments` □□□□□□□Sequential□□□□□□□
      - □ `nn.Sequential` □□□□ `segments` □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□
      - □□□□□□□□□ `segments` □□□□□□□□□□□□□□□

□□

```
from torch.utils.checkpoint import checkpoint_sequential

model = nn.Sequential(
    nn.Linear(10, 10),
    nn.ReLU(),
    nn.Linear(10, 10),
    nn.ReLU(),
```

```
    nn.Linear(10, 10)
)

input_tensor = torch.randn(2, 10, requires_grad=True)
## □□□□ chunk □□2□□□□ forward □□ checkpoint
output = checkpoint_sequential(model, chunks=2, input=input_tensor)
```

□□□□

□□□□

```
## □□□□□□
accumulation_steps = 4
## □□□□
optimizer.zero_grad()

for i ,(inputs,labels) in enumerate(datalaoder):
    outputs = model(inputs)
    # □□□□
    loss = criterion(outputs,labels)
    # □□□□□□□□□□□□□□□□□□□□□□□
    loss = loss / accumulation_steps
    # □□□□
    loss.backward() # □□□□□□□□□□□□□□□□

    if (i + 1) % accumulation_steps == 0:
        # □□□□
        optimizer.step()
        # □□□□
        optimizer.zero_grad()
```

□□□

- □□□□□□□□$batch_size$=36□

$$ \nabla L_{true} = \frac{1}{32}\sum_{i = 1}^{32}\nabla loss_i $$

- □□□□□□4□□$mini_batch$ = 8□

$$ \nabla L_{accu} = \sum_{j=1}^4(\frac{1}{4}\cdot \frac{1}{8}\sum_{k=1}^8\nabla loss_{jk}) = \frac{1}{32}\sum_{i=1}^{32}\nabla loss_i $$

□□□□

□ `torch.nn.utils` □□□□□□□□□□□□□□□□□□□□

□□□□□□□□□□□□□□□

□□□□□ `clip_grad_value_` □

- □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□
- `torch.nn.utils.clip_grad_value_`
    - `model.parameters()` □□□□□

- `clip_value` □□□□□□

□□□□□□ `clip_grad_norm_` □

- □□□□□□**L2**□□□□□□□□□□□□□□□□□□□□□□□□
- `torch.nn.utils.clip_grad_norm_`
  - `model.parameters()` □□□□□
  - `max_norm` □□□□□L2□□

□□□□

pytorch□□□□□□□□□□□□

```
class EarlyStopping:
    def __init__(self,patience=3):
        self.patience = patience
        self.best_loss = float('inf')
        self.counter = 0
    # □□□□□□□□□□□□□□
    def __call__(self,val_loss):
        if val_loss < self.best_loss:
            self.best_loss = val_loss
            self.counter = 0
            return False
        else:
            self.counter += 1
            return self.counter >= self.patience
```

□□□□□

□□

□□□□□□□

□□□□□□□□□□□□□ `int8` □□□□□ `float32`

- □ `torch.quantization` □□ `from torch.quantization import quantize_dynamic`
  - `model` □□□□□
  - `{nn.Linear,nn.LSTM}` □□□□□□□□□
  - `dtype` □□□□□□□□□

□□□□□□□□□□□□□ `int8`

- □□□
  1. □□□□□□
  2. □□□□□ `fuse_modules` □
     - `torch.quantization.fuse_modules`
  3. □□□□□□ `qconfig` □
     - `torch.quantization.get_default_qconfig`
       - `fbgemm` □□x86□ `qnnpack` □□ARM

```
model.qconfig = torch.quantization.get_default_qconfig('fbgemm')
```

4. □□□□□ `prepare` □
    - `torch.quantization.prepare()`
        - `inplace` □□□□□□□□□
5. □□□□□□□□□□□□ forward□
6. □□□□□□□□□ `convert` □
    - `torch.quantization.convert`

□□□□

```
model = MyModel()
model.eval()

## 1. □□□□
model_fused = torch.quantization.fuse_modules(model, [['conv', 'relu']])

## 2. □□qconfig
model_fused.qconfig = torch.quantization.get_default_qconfig('fbgemm')

## 3. □□□□
torch.quantization.prepare(model_fused, inplace=True)

## 4. □□
with torch.no_grad():
    for data, _ in calibration_loader:
        model_fused(data)

## 5. □□
torch.quantization.convert(model_fused, inplace=True)
```

□□□□□□□**(QAT)**□□□□□□□□“□□□□”□□□□□□□ `FakeQuantize` □□□□□□□□□□□□□□□

- □□
    1. □□□□□
        - `torch.quantization.fuse_modules`
    2. □□ `qconfig` □□□ QAT □□□
        - `torch.quantization.get_default_qat_qconfig`
    3. □□ `prepare_qat`
        - `torch.quantization.prepare_qat()`
    4. □□□□□ epoch
    5. `convert` □□□□□□□
        - `torch.quantization.convert()`

□□□□

```python
model.train()
model_fused = torch.quantization.fuse_modules(model, [['conv', 'relu']])

## 1. □□□ QAT □□
model_fused.qconfig = torch.quantization.get_default_qat_qconfig('fbgemm')

## 2. □□ QAT
torch.quantization.prepare_qat(model_fused, inplace=True)

## 3. □□□□
for epoch in range(num_epochs):
    for data, target in train_loader:
        output = model_fused(data)
        loss = loss_fn(output, target)
        loss.backward()
        optimizer.step()

## 4. □□□□□□□□□□
model_fused.eval()
torch.quantization.convert(model_fused, inplace=True)
```

□□

□ `torch.nn.utils.prune` □

□□□□□

- `prune.l1_unstructured` □□**L1**□□□□□□□□□□□□
  - □□□□□□□□□□□□□□□□□□□□**L1**□□
  - □□□□□□□□□□□□□□□□□□□□**0**
- `prune.random_unstructured` □□□□□□□□□□□□□
  - `module` □□□□□□□
  - `name` □□□□□□□□□□□□□□□□
  - `amount` □
    - □□□ `float` □□□□□□□□□
    - □□□ `int` □□□□□□□□□□□
- `prune.ln_structured` □□□□/□□
  - □□□□□
  - `amount` □
    - □□□ `float` □□□□□□□□□
    - □□□ `int` □□□□□□□□□□□
  - `n` □L-n□□
  - `dim` □□□□□□□
    - `0` □□□□
    - `1` □□□□
- `prune.remove` □□□□□□□□□□□□
  - `module` □□□□□□□
  - `name` □□□□□□□□□□□□□□□□

□□□□□□□□□□□□□□□□□□□□□□□□□□□□□mask□0□□□□□□□□□□□□□□

□□□□□

□□□tensorboard

□ `torch.utils` □□□□ `from torch.ultis.tensorboard import SummaryWriter`

- □□□□□□□□□□□□□□□□□□□□□□□□□tensorboard□□,□□□□□□□□□□

```python
from torch.utils.tensorboard import SummaryWriter
## □□ SummaryWriter □□□□□□□

writer = SummaryWriter('C:/study/pytorch/logs')

## □□ SummaryWriter

writer.close()
```

1. □□□□
   - `writer.add_scalar()`
     - `tag` □□□□□
       - □□ `train/loss` □ `eval/loss` □□□□□□□□□
     - `scalar_value` □□□□□□□□□□□□y□
     - `global_step` □□□□□□□□□□□□x□

2. □□□□□□□□□□□□□□□
   - `writer.add_graph()`
     - □□
     - □□□□

3. □□□□□□□□□□□
   - `writer.add_hparams()`
     - □□□□□□□□□□□
     - □□□□□□□□□□

tqdm

□□□□□□□□□□

`from tqdm import tqdm`

`tqdm()`

- **□□□□□□□□□Dataloader□□**
- `desc` □□□□□□□□□□
- `total` □□□□□
- `leave` □□□□□□□□□□□

- `unit` □□□□□`it/s`□□`batch/s`□
- `ncols` □□□□□

`□□.set_postfix()` □□□□□□□□□□□□

- □□□ `train_bar.set_postfix(loss=f"{loss.item():.4f}")`
- □□□

  ```
  train_bar.set_postfix({
          'loss': f"{loss.item():.4f}",
          'acc': f"{acc.item():.2%}"  # □□□□□
      })
  ```

□□□□□□

- □□□□□

`torch.save(model.state_dict(), './model_weights.pth')`

- □□□□□□□□

`model.load_state_dict(torch.load('./model_weights.pth'))`

- □□□□□□□□□

`torch.save(model,path)`

- □□□

`model = torch.load(path)`

□□□□

□□□□

- `from torch.nn.parallel import DistributedDataParallel as DDP`
- `from torch.utils.data.distributed import DistributedSampler`
- `import torch.distributed as dist`
- `import torch.multiprocessing as mp`

□□□□

1. □□□□□

□□□□□□□□□□□□□GPU

2. □□□□

□□□□□□□□□□□□□

3. □ `DDP` □□□□□□ `DistributedSampler` □□□
4. □□ `torch.multiprocessing.spawn` □□□□

□□□

- `loss.backward()` □□□□□□□□□□□□□□□□□ `.grad` □□□□□□□□□□□□□□□□□□□□□□□□□□□

- □□□□□ `optimizer.step()` □□□□□□□
- □□□□□ `optimizer.zero_grad()` □□□□□□□□
- □□□□□□□□□□□□□□□□□□□□□□□□□□□□□

* `torch.no_grad()` □□□□□□□□□□□□
* `.requires_grad_` □□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□□
* `loss.item()` □□□□□□□□□□ **Python**□□
* `optim.zero_grad()` □□□□□
* `optimizer.step()` □□□□□
* `loss.backward()` □□□□□
* `self.training` □ `nn.Module` □□□□□□□□□□□□□□□□□□□□□□□□□□ `True` □□□□□□□ `False` □
* □□□□□□□□□□□□□□
* □□□□□□□□□□□□□□□□□□□□□□
* `.permute()` □□□□□□□□□□□□□□□□□□□□ `.reshape()` □□□□□□□□□□□□□□□□□□□